

energy_py

lessons learnt building an energy reinforcement
learning library



4 years energy engineer @ **ENGIE**

graduate & 2 years teaching reinforcement
learning @ **Data Science Retreat**

data scientist @ **Tempus Energy**

blog at www.adgefficiency.com



agenda

motivation

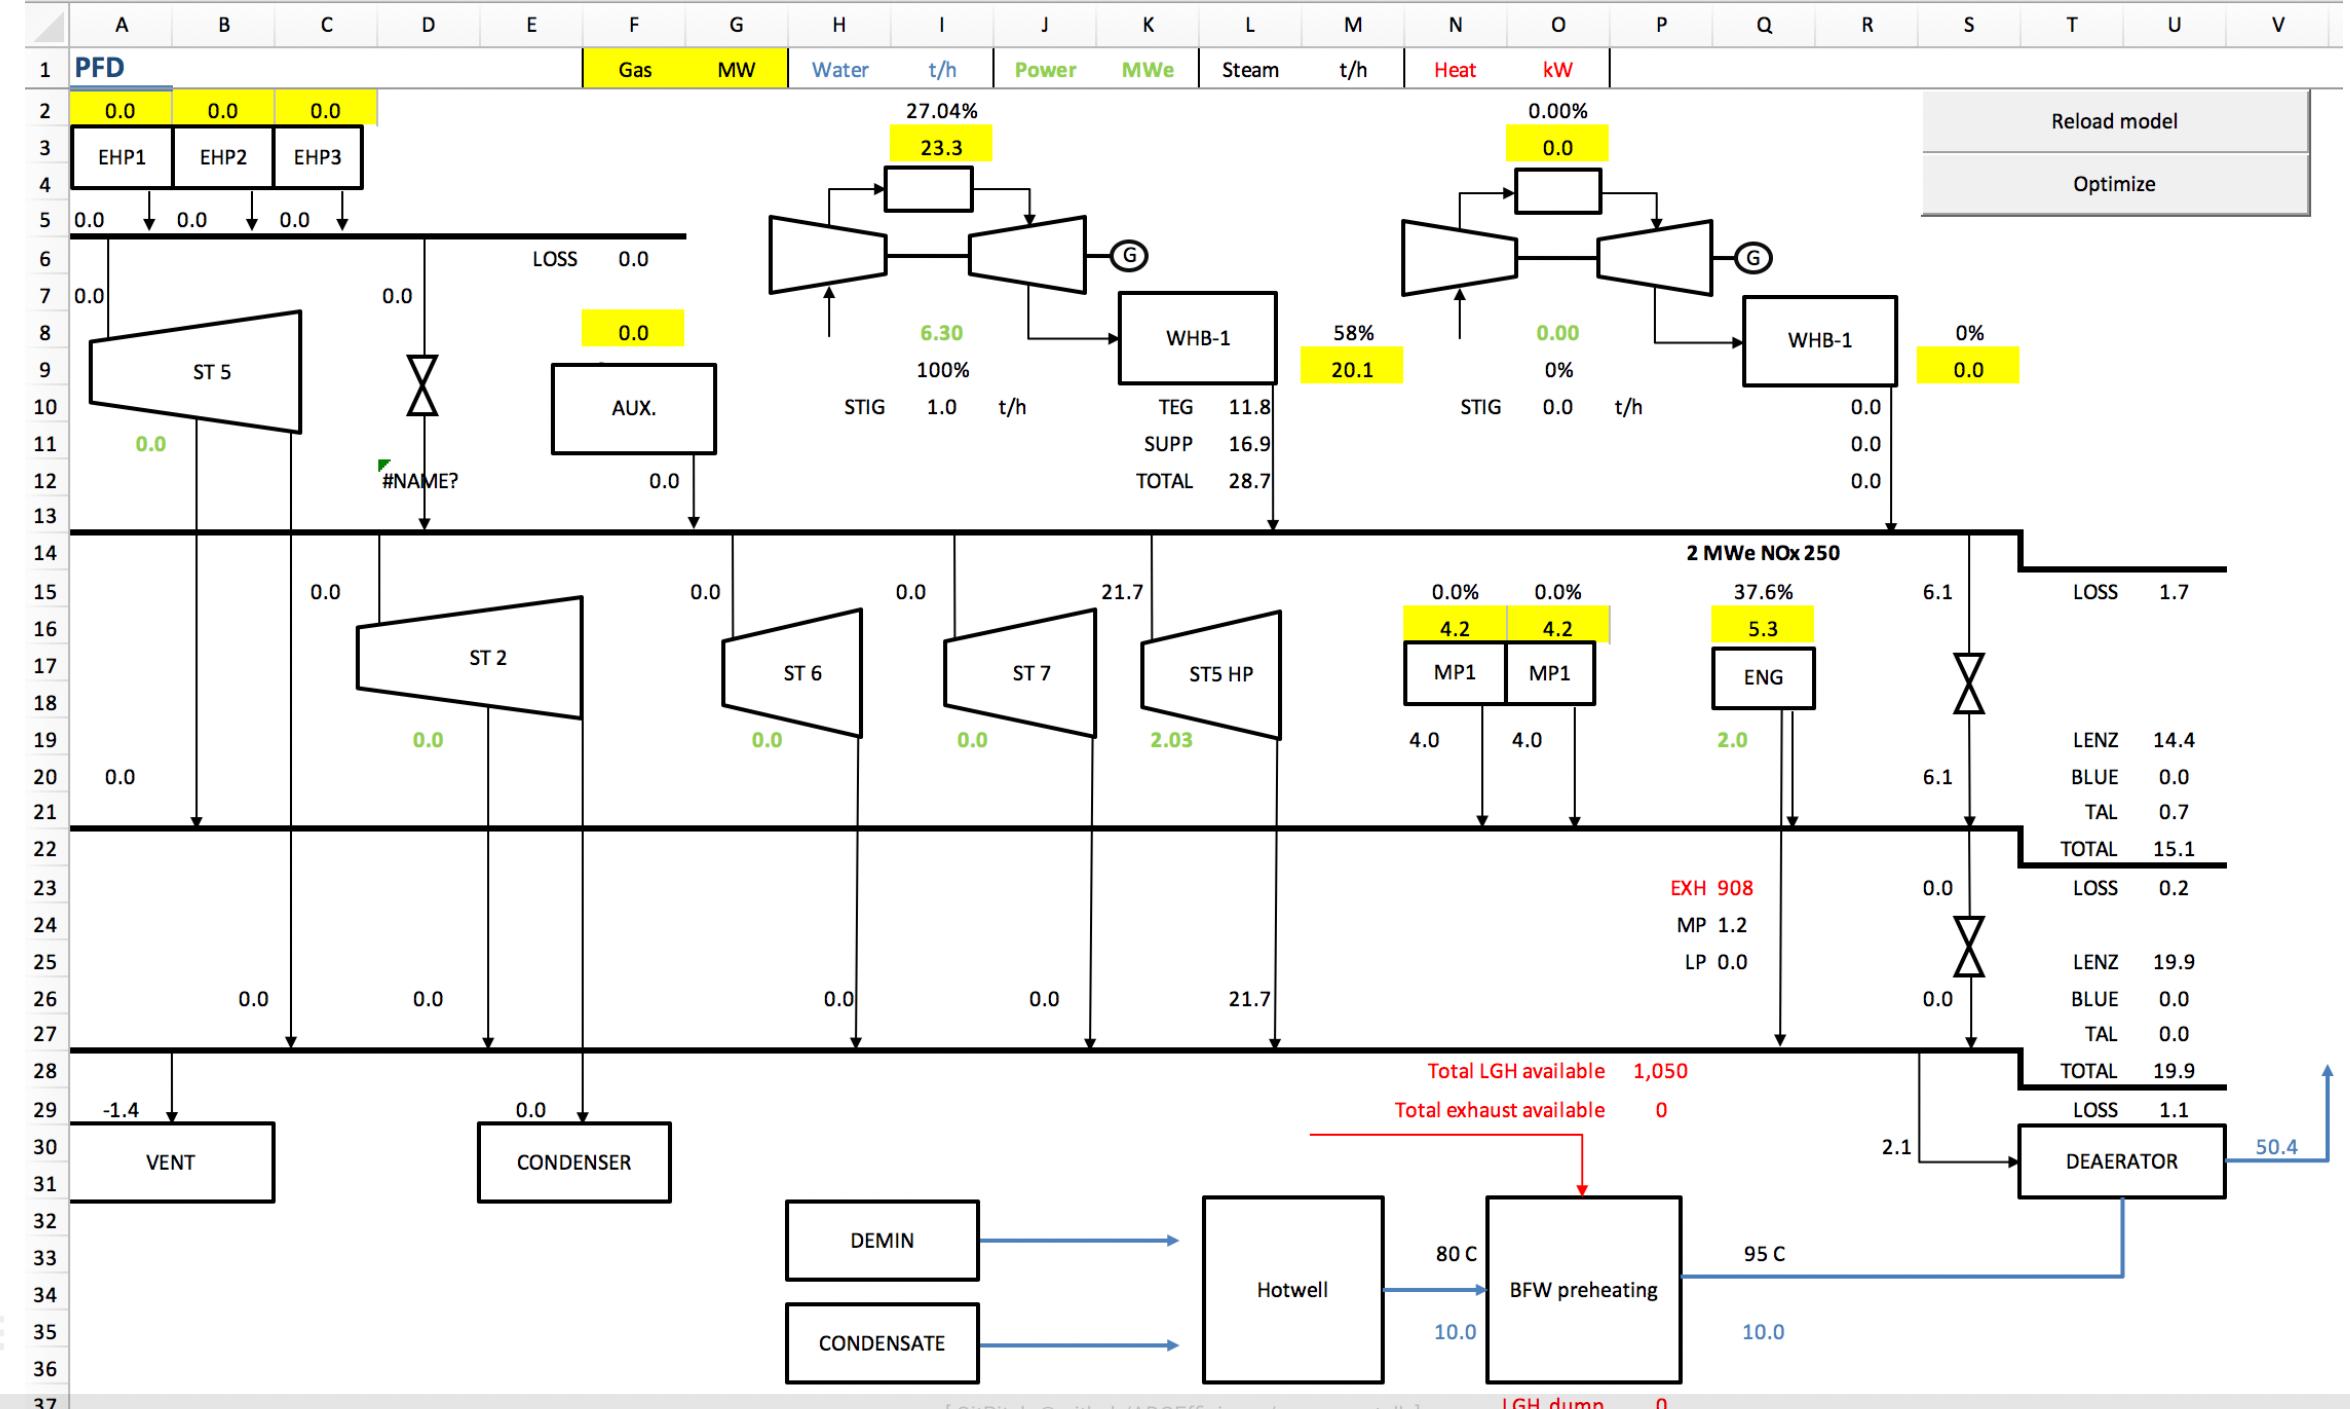
API + performance

lessons

next steps







computation to help solve the **control** problem

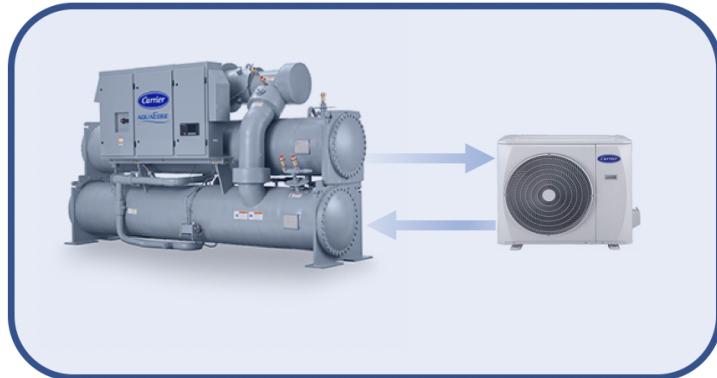
computation to help solve the **climate** problem



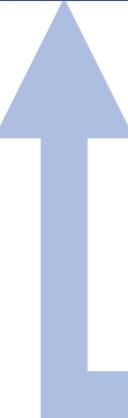
price response flexible demand

aka the lazy taxi driver

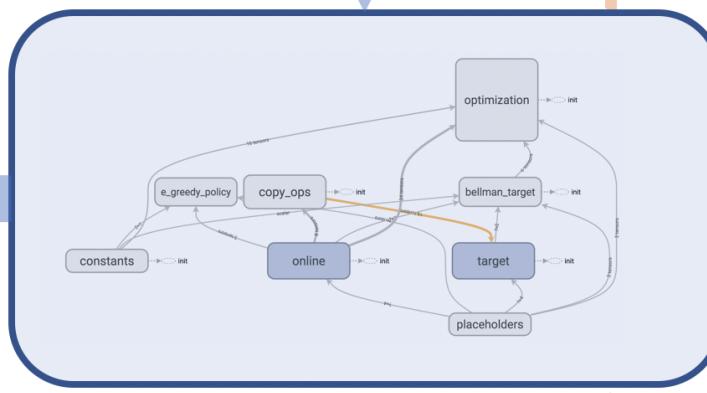
environment



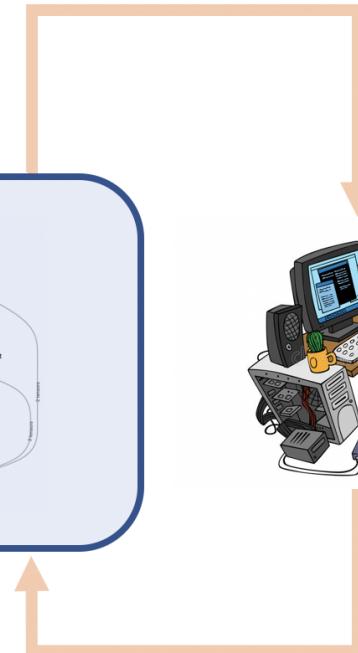
state = cooling demands, electricity prices
reward = electricity consumption * electricity price



action = temperature setpoint

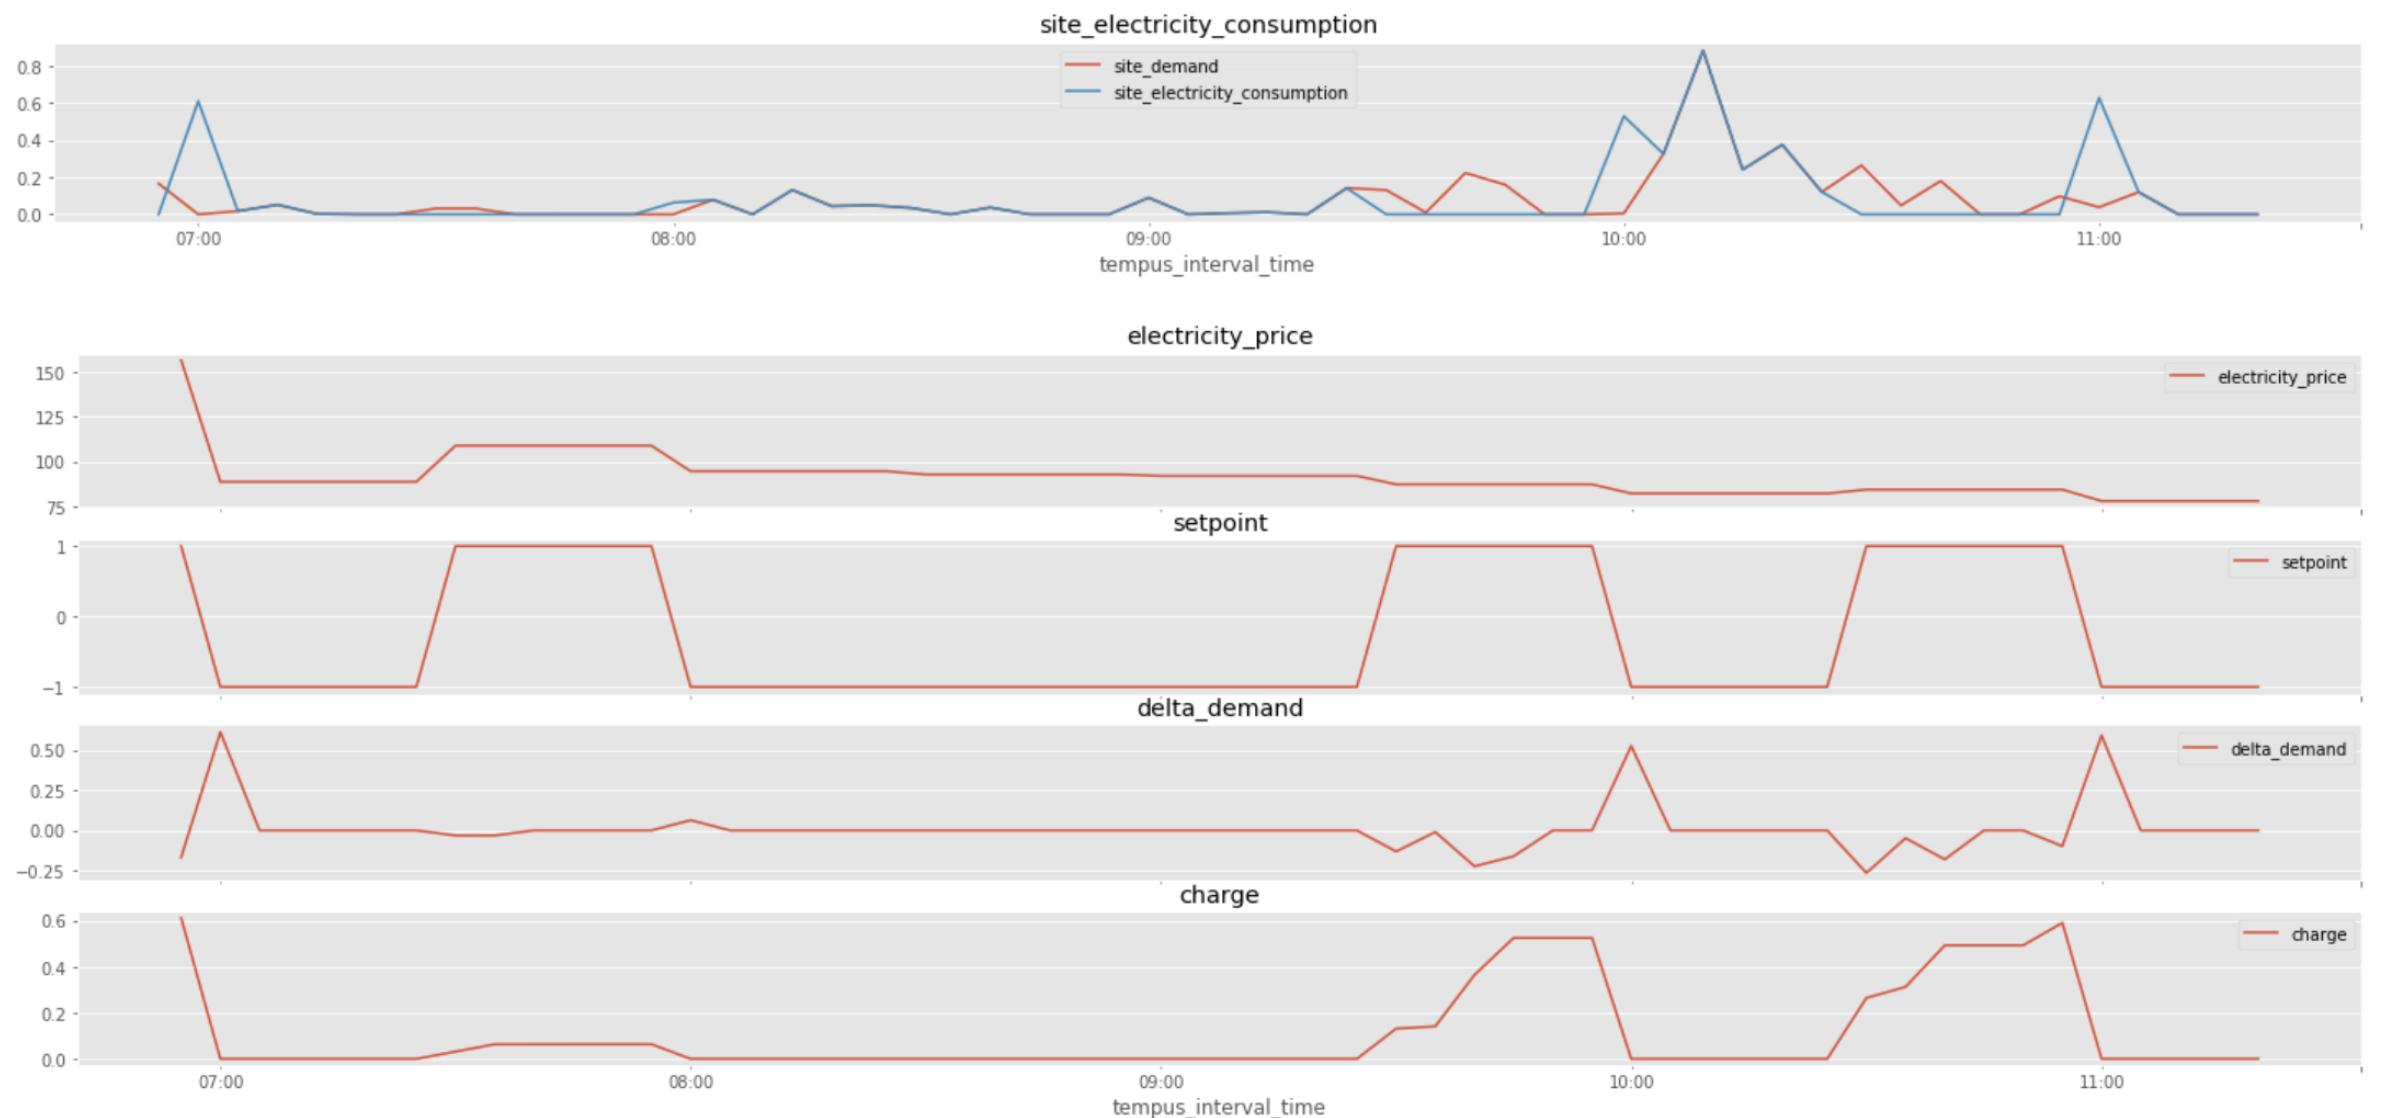


agent



hyperparameter
torturer





 [ADGEfficiency / energy_py](#)

[Watch](#) 8 | [Star](#) 62 | [Fork](#) 15

[Code](#) | [Issues 14](#) | [Pull requests 1](#) | [Projects 0](#) | [Wiki](#) | [Insights](#) | [Settings](#)

reinforcement learning agents and environments for energy systems [Edit](#)

[reinforcement-learning](#) [energy](#) [reinforcement-learning-agents](#) [energy-systems](#) [Manage topics](#)

 702 commits  2 branches  0 releases  5 contributors  MIT

Branch: master ▾ [New pull request](#) [Create new file](#) [Upload files](#) [Find file](#) [Clone or download](#) ▾

	agents	aws sync
	common	array memory remember forms np.array for all experience elements befo...
	envs	readme update
	experiments	env info refactor
	tests	conv net working with 2048
	_init__.py	tests working with network register

DQN + naive agents

energy environments + wrappers around gym

tools for experimentation



High level API

```
$ cd energy_py/experiments  
$ python experiment.py example dqn  
$ tensorboard --logdir='./results/example/tensorboard'
```



Low level API

```
import energy_py

with tf.Session() as sess:
    env = energy_py.make_env(
        env_id='battery',
        episode_length=288,
        dataset='example'
    )

    agent = energy_py.make_agent(
        sess=sess,
        agent_id='dqn',
        env=env,
        total_steps=1000000,
        batch_size=1024,
    )
```

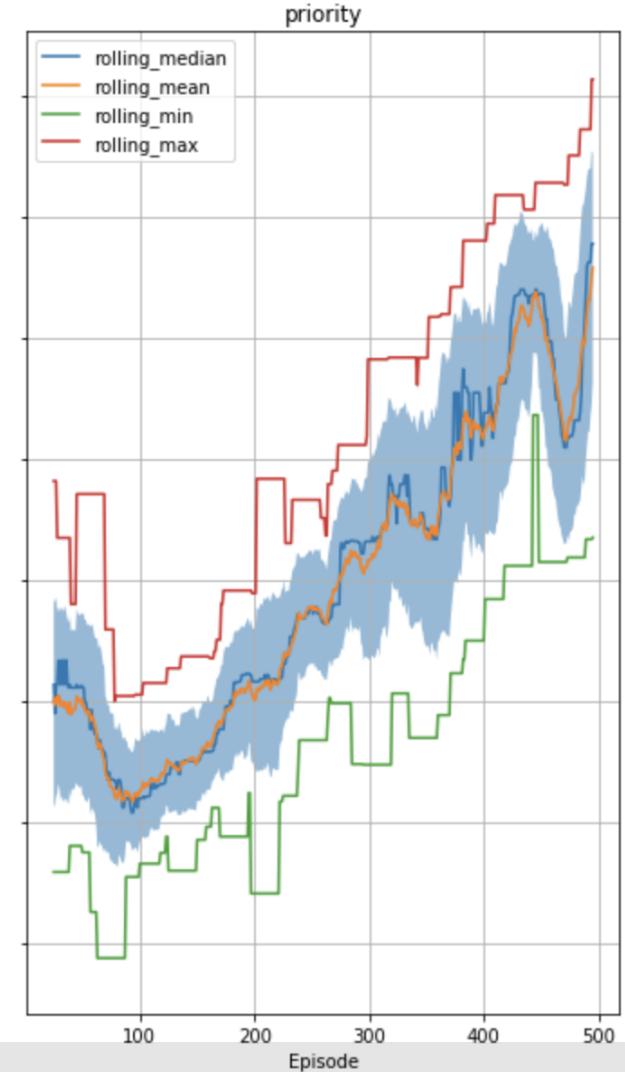
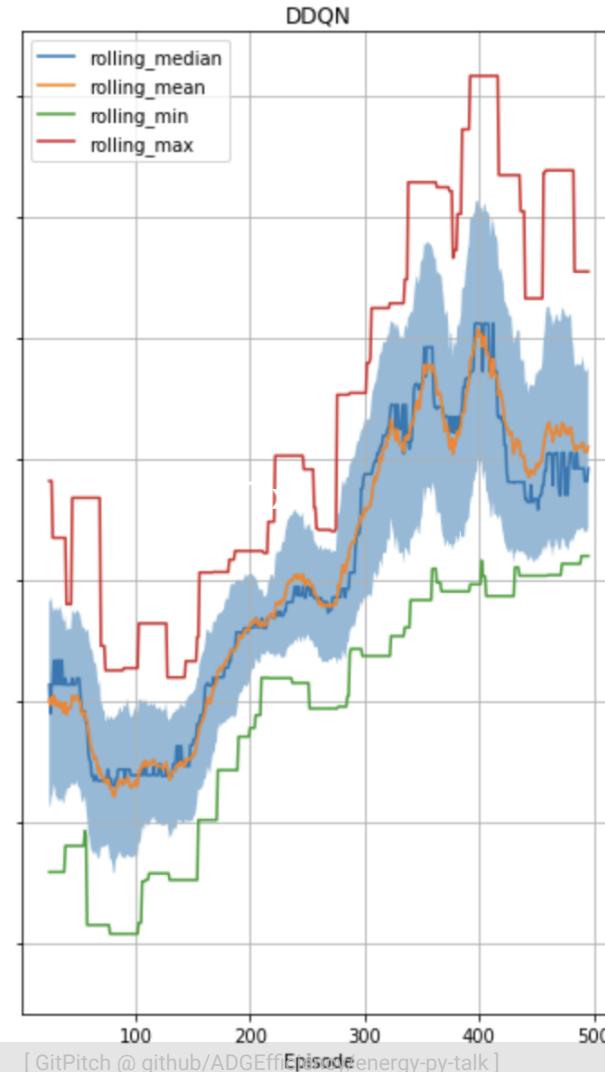
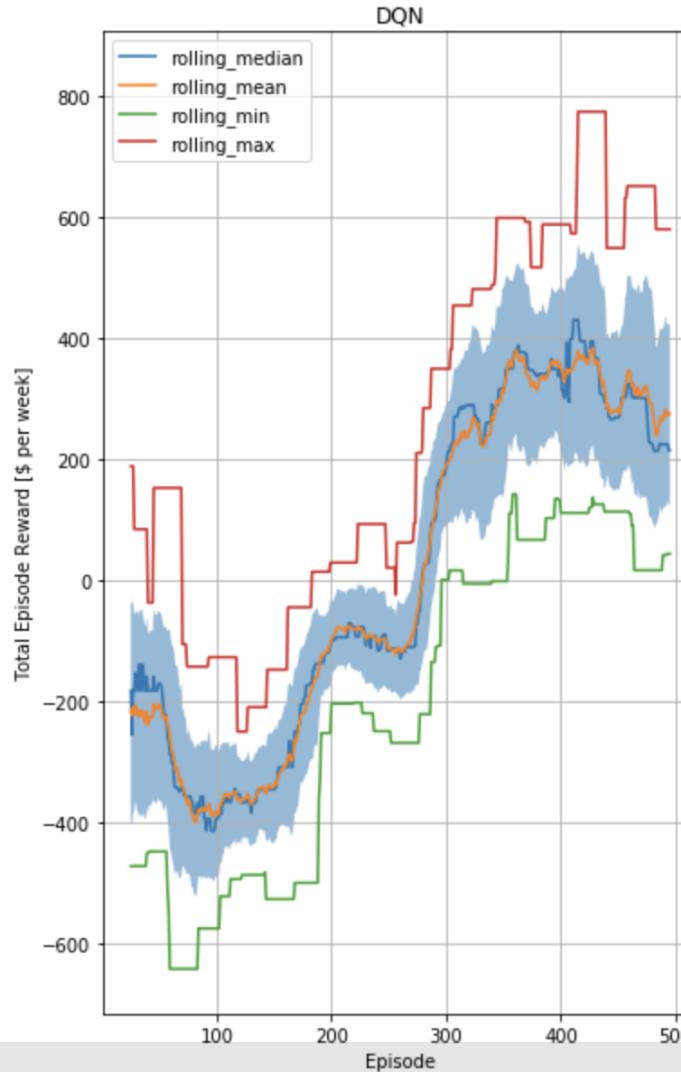


fun with cartpole



battery storage

Average over 5 random seeds. Window length of 25 episodes. Shaded area shows +/- the rolling standard deviation.



details - space design



```
#  create an action space of two dimensions
action_space = GlobalSpace('action').from_spaces(
    [ContinuousSpace(0, 100), DiscreteSpace(3)],
    ['acceleration', 'gear']
)
#  randomly sample an action
action = action_space.sample()
#  check the action is valid
assert action_space.contains(action)
#  discretize the space
discrete_spaces = action_space.discretize(20)
#  randomly sample a discrete action
action = action_space.sample_discrete()
```



```
# load a state or observation space from a dataset
state_space = GlobalSpace('state').from_dataset('example')

# we can sample an episode from the state
episode = state_space.sample_episode(start=0, end=100)

# sample from the current episode by calling the space
state = state_space(steps=0)
```



rl today

hard

unstable

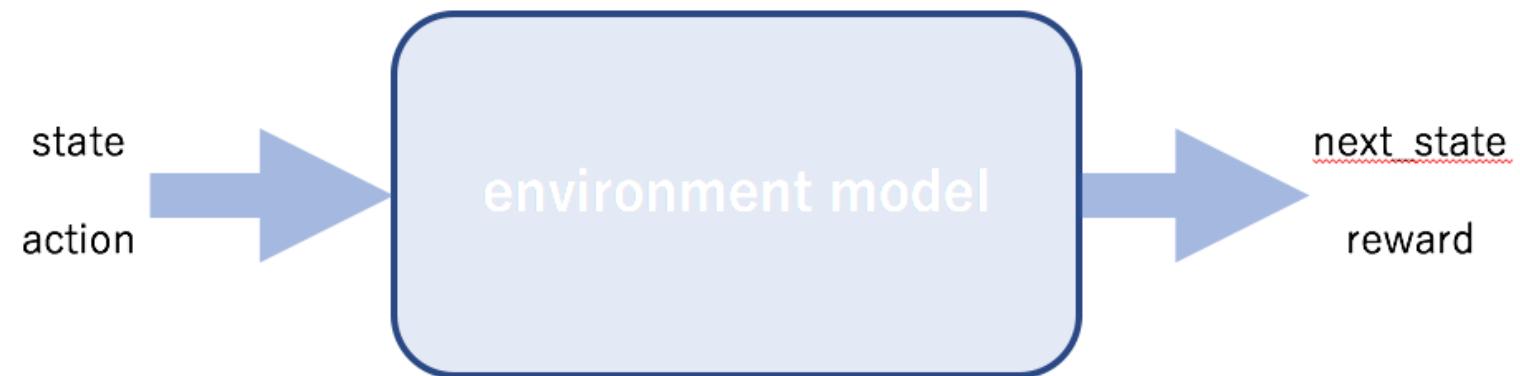
sample inefficient

exciting



the environment model problem / opportunity





In a similar vein, you can easily outperform DQN in Atari with off-the-shelf Monte Carlo Tree Search. Here are baseline numbers from [Guo et al, NIPS 2014](#). They compare the scores of a trained DQN to the scores of a UCT agent (where UCT is the standard version of MCTS used today.)

Agent	<i>B.Rider</i>	<i>Breakout</i>	<i>Enduro</i>	<i>Pong</i>	<i>Q*bert</i>	<i>Seaquest</i>	<i>S.Invaders</i>
DQN	4092	168	470	20	1952	1705	581
<i>-best</i>	5184	225	661	21	4500	1740	1075

Agent	<i>B.Rider</i>	<i>Breakout</i>	<i>Enduro</i>	<i>Pong</i>	<i>Q*bert</i>	<i>Seaquest</i>	<i>S.Invaders</i>
UCT	7233	406	788	21	18850	3257	2354

```
@misc{rlblogpost,
    title={Deep Reinforcement Learning Doesn't Work Yet},
    author={Irpan, Alex},
    howpublished={\url{https://www.alexirpan.com/2018/02/14/rl-hard.html}},
    year={2018}
}
```



backwards induction

```
previous_state_payoffs = {state: 0 for state in model.states}

for step in steps[::-1]:
    viable_transitions = [get_viable_transitions(step, next_state, model)
                          for next_state in model.states]

    new_state_payoffs = defaultdict(list)

    for t in viable_transitions:
        payoff = t.reward + previous_state_payoffs[t.next_state]
        new_state_payoffs[t.state].append(payoff)

    for k, v in new_state_payoffs.items():
        new_state_payoffs[k] = np.max(v)

    previous_state_payoffs = new_state_payoffs
```



synthetic data - aka poor mans GANS



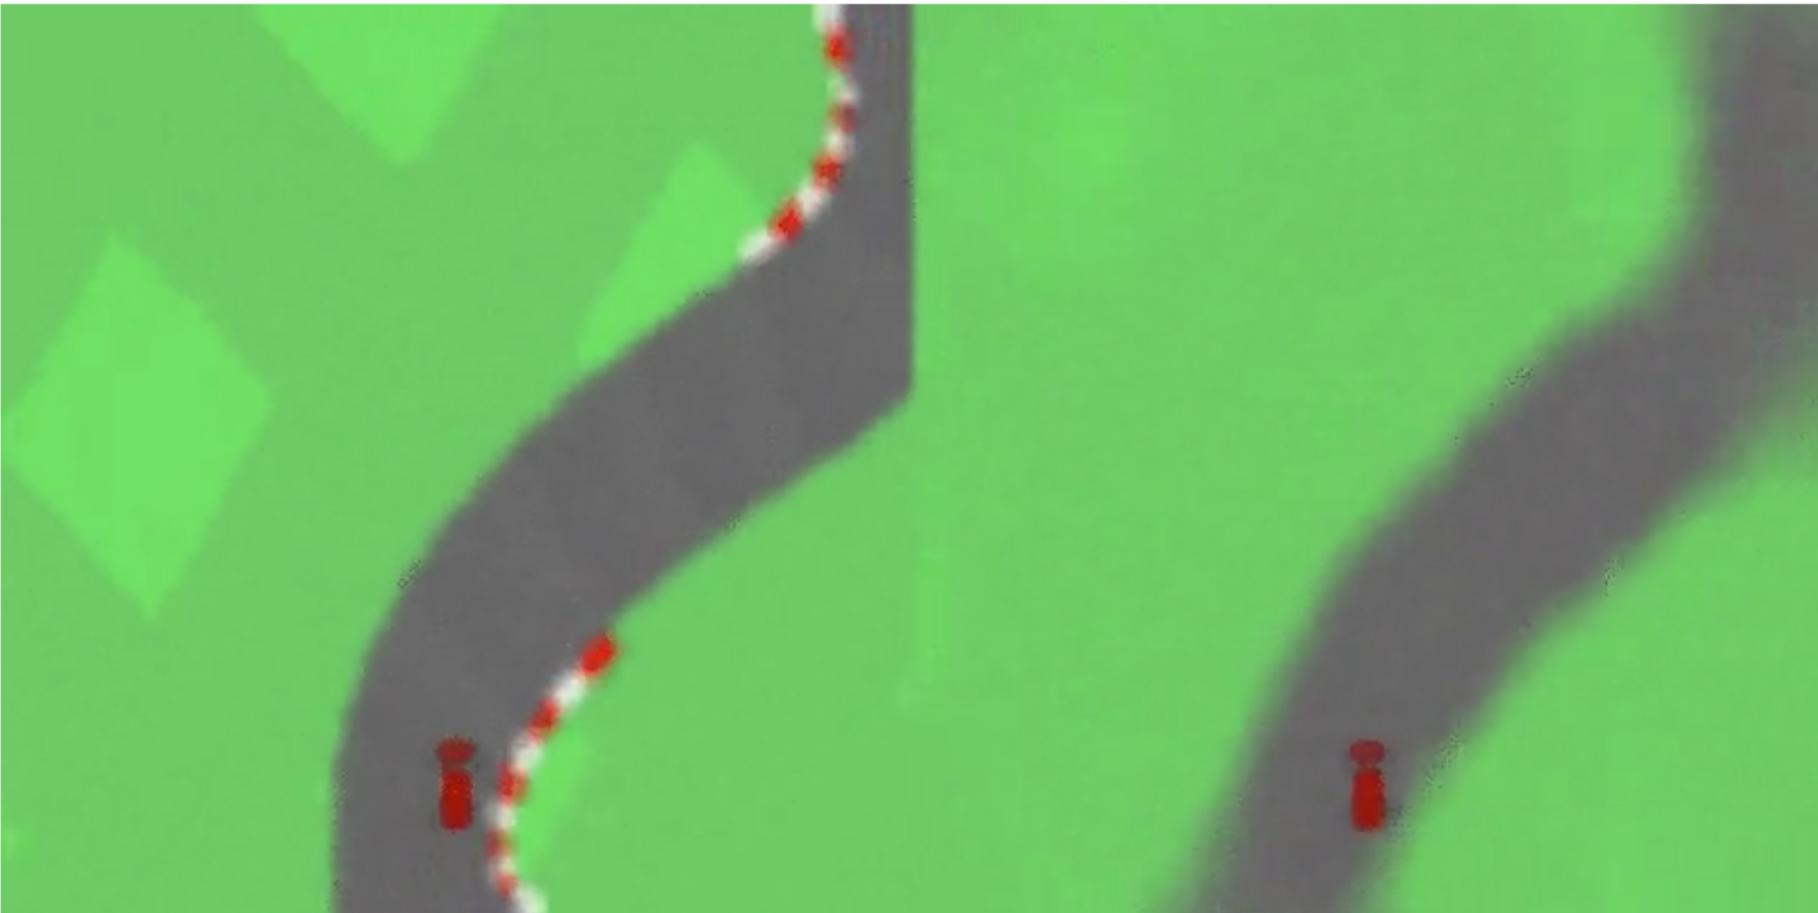
World Models

David Ha, Jürgen Schmidhuber

(submitted on 27 Mar 2018 ([v1](#)), last revised 9 May 2018 (this version, v4))

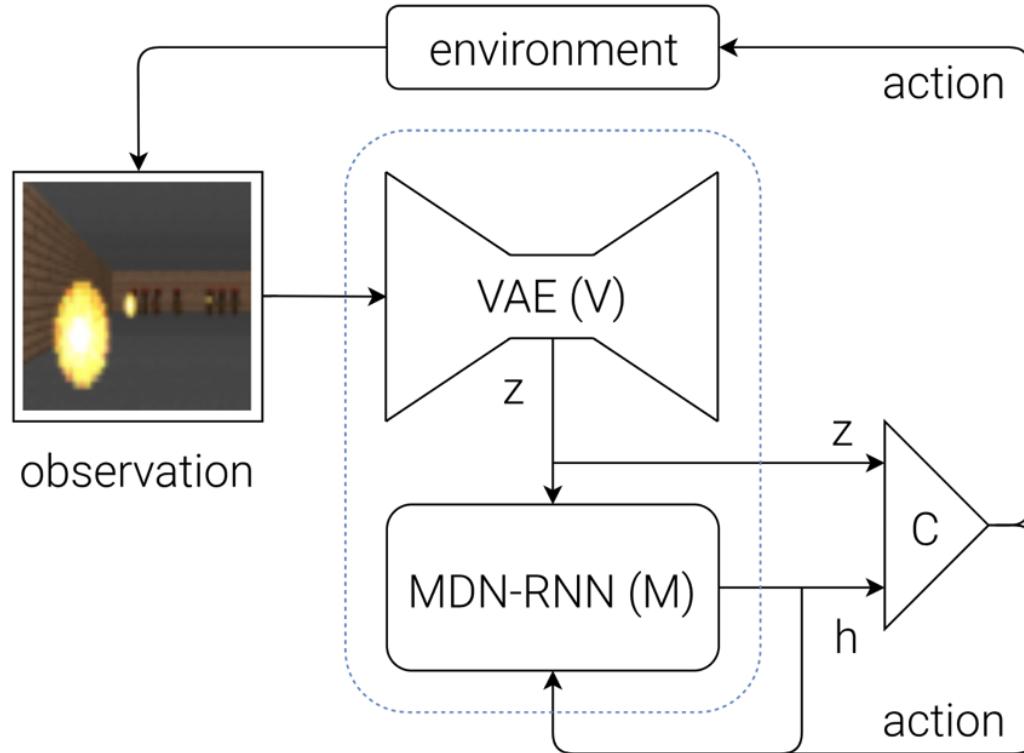
We explore building generative neural network models of popular reinforcement learning environments. Our world model can be trained quickly in an unsupervised manner to learn a compressed spatial and temporal representation of the environment. By using features extracted from the world model as inputs to an agent, we can train a very compact and simple policy that can solve the required task. We can even train our agent entirely inside of its own hallucinated dream generated by its world model, and transfer this policy back into the actual environment.



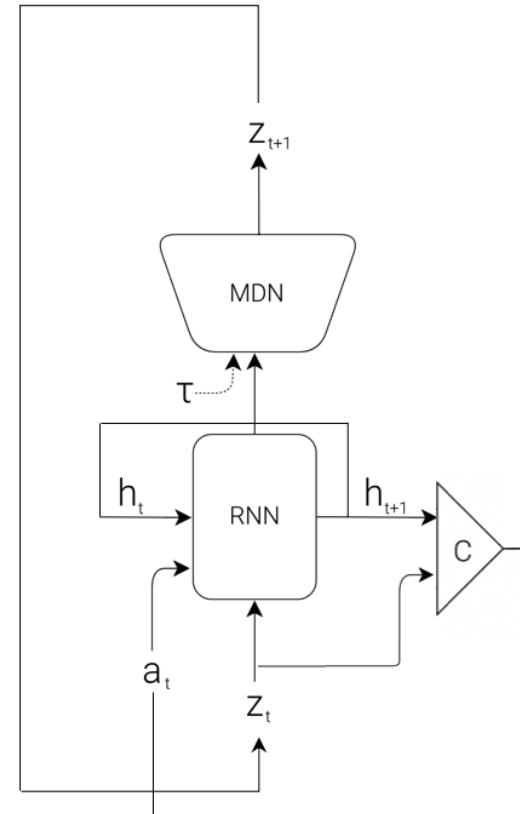


Actual observations from the environment.

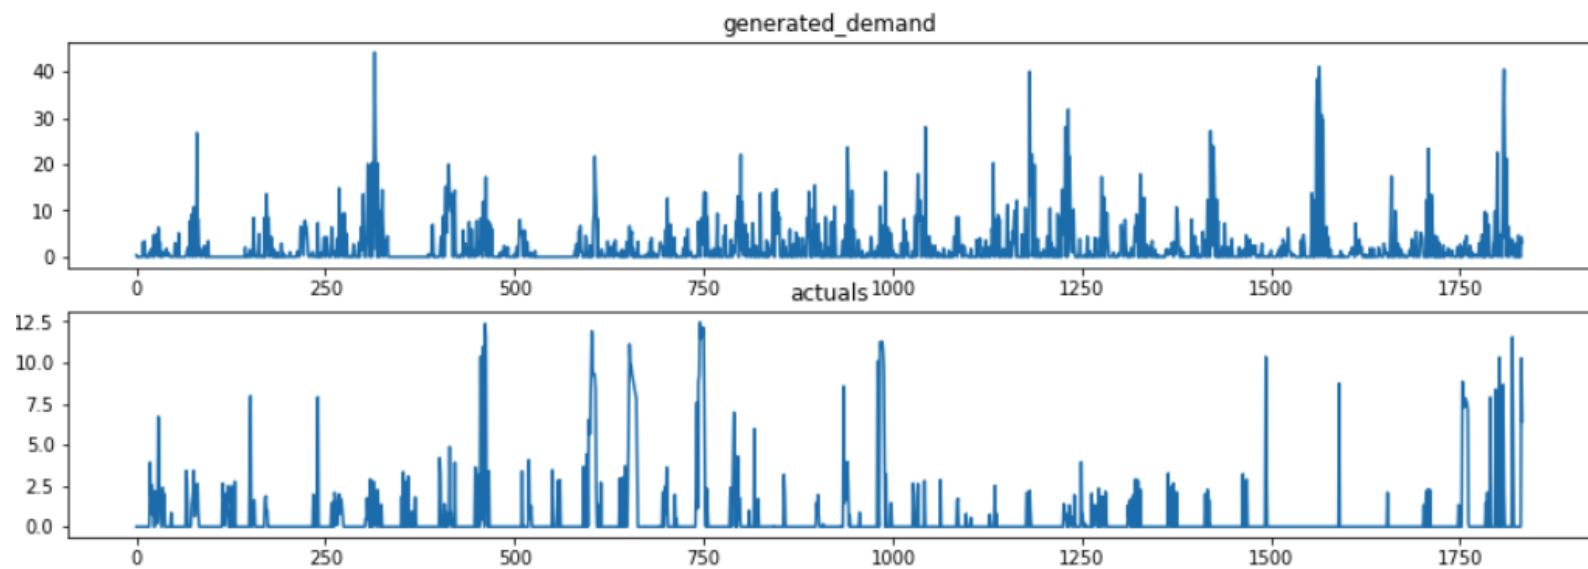
What gets encoded into z_t .



Learning using the real environment



Learning inside a dream



long term work

a fourth complete rebuild

wrapping other environments

model based methods

integrating with google/dopamine



Dopamine



Dopamine is a research framework for fast prototyping of reinforcement learning algorithms. It aims to fill the need for a small, easily grokked codebase in which users can freely experiment with wild ideas (speculative research).

Our design principles are:

- *Easy experimentation*: Make it easy for new users to run benchmark experiments.
- *Flexible development*: Make it easy for new users to try out research ideas.
- *Compact and reliable*: Provide implementations for a few, battle-tested algorithms.
- *Reproducible*: Facilitate reproducibility in results.

In the spirit of these principles, this first version focuses on supporting the state-of-the-art, single-GPU *Rainbow* agent ([Hessel et al., 2018](#)) applied to Atari 2600 game-playing ([Bellemare et al., 2013](#)). Specifically, our Rainbow agent implements the three components identified as most important by [Hessel et al.](#):

- n-step Bellman updates (see e.g. [Mnih et al., 2016](#))
- Prioritized experience replay ([Schaul et al., 2015](#))
- Distributional reinforcement learning ([C51; Bellemare et al., 2017](#))



thank you

